

HW9_Hyunwoo

March 20, 2019

1 HW 9

2 Name: Hyunwoo Roh

3 Question1

Neural network horse race.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn import preprocessing
from sklearn.linear_model import LogisticRegression
from sklearn import metrics
from sklearn.model_selection import RandomizedSearchCV, GridSearchCV
from sklearn.ensemble import RandomForestRegressor
from sklearn.svm import SVC
from scipy.stats import randint as sp_randint
from scipy.stats import uniform as sp_uniform
from sklearn.neural_network import MLPClassifier
```

```
In [4]: # read in data
drink=pd.read_csv("data\strongdrink.txt")
drink.head()
```

```
Out[4]:
```

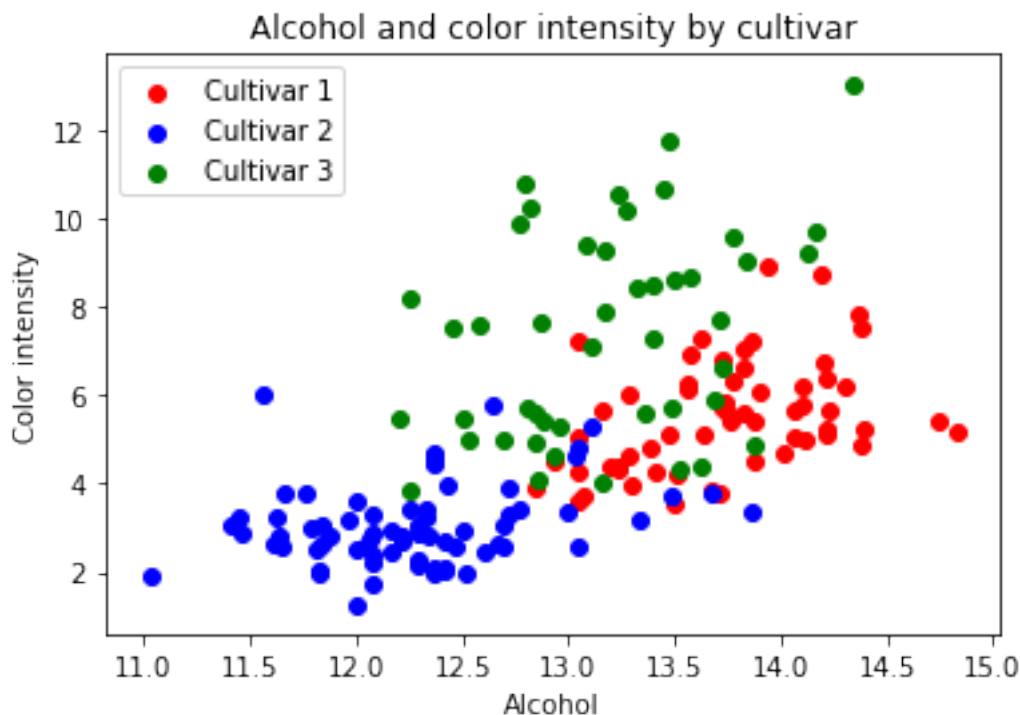
	cultivar	alco	malic	ash	alk	magn	tot_phen	flav	nonfl_phen	\
0	1	14.23	1.71	2.43	15.6	127	2.80	3.06	0.28	
1	1	13.20	1.78	2.14	11.2	100	2.65	2.76	0.26	
2	1	13.16	2.36	2.67	18.6	101	2.80	3.24	0.30	
3	1	14.37	1.95	2.50	16.8	113	3.85	3.49	0.24	
4	1	13.24	2.59	2.87	21.0	118	2.80	2.69	0.39	

	proanth	color_int	hue	OD280rat	proline
0	2.29	5.64	1.04	3.92	1065
1	1.28	4.38	1.05	3.40	1050
2	2.81	5.68	1.03	3.17	1185
3	2.18	7.80	0.86	3.45	1480
4	1.82	4.32	1.04	2.93	735

3.0.1 Create a scatterplot of the data where the x-variable is alcohol and y variable is color intensity

```
In [10]: is_1 = drink.cultivar == 1
         is_2 = drink.cultivar == 2
         is_3 = drink.cultivar == 3

         plt.scatter(drink[is_1].alco, drink[is_1].color_int, color='r', label="Cultivar 1")
         plt.scatter(drink[is_2].alco, drink[is_2].color_int, color='b', label="Cultivar 2")
         plt.scatter(drink[is_3].alco, drink[is_3].color_int, color='g', label="Cultivar 3")
         plt.xlabel("Alcohol")
         plt.ylabel("Color intensity")
         plt.title("Alcohol and color intensity by cultivar")
         plt.legend()
         plt.show()
```



3.1 (b) Use sk.learn to fit a multinomial logistic model of cultivar on given features.

```
In [11]: x_variable_names = ['alco', 'malic', 'tot_phen', 'color_int']
         y_name = ['cultivar']
         Xvars = drink[x_variable_names].values
         yvals = drink[y_name].values

In [13]: clf_mlog = LogisticRegression(solver='newton-cg', multi_class='multinomial').fit(Xvars,
         param_dist1 = {'penalty': ['l1', 'l2'], 'C': sp_uniform(0.1, 10.0)})
```

```

random_search1 = RandomizedSearchCV(LogisticRegression().fit(Xvars, yvals),
                                     param_distributions=param_dist1,
                                     n_iter=200, n_jobs=-1, cv=5, random_state=25,
                                     scoring='neg_mean_squared_error')

random_search1.fit(Xvars,yvals)
print('RandBestEstimator1=', random_search1.best_estimator_)
print('RandBestParams1=', random_search1.best_params_)
print('RandBestScore1=', -random_search1.best_score_)

RandBestEstimator1= LogisticRegression(C=2.665871587495725, class_weight=None, dual=False,
fit_intercept=True, intercept_scaling=1, max_iter=100,
multi_class='warn', n_jobs=None, penalty='l1', random_state=None,
solver='warn', tol=0.0001, verbose=0, warm_start=False)
RandBestParams1= {'C': 2.665871587495725, 'penalty': 'l1'}
RandBestScore1= 0.11931818181818182

```

3.2 (c) Use sklearn randomforest to fit a random forest model of cultivar on the same four features used in part (b)

```

In [14]: rfc = RandomForestRegressor(bootstrap=True,oob_score=True, random_state=25).fit(Xvars

In [16]: param_dist2 = {'n_estimators':sp_randint(10,200),
                        'max_depth': sp_randint(2,4),
                        'min_samples_split': sp_randint(2, 20),
                        'min_samples_leaf': sp_randint(2, 20),
                        'max_features': sp_randint(1, 4)}

random_search2 = RandomizedSearchCV(rfc, param_distributions=param_dist2,
                                     n_iter=200, n_jobs=-1, cv=5, random_state=25,
                                     scoring='neg_mean_squared_error')

random_search2.fit(Xvars,yvals)
print('RandBestEstimator2=', random_search2.best_estimator_)
print('RandBestParams2=', random_search2.best_params_)
print('RandBestScore2=', -random_search2.best_score_)

RandBestEstimator2= RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=3,
max_features=3, max_leaf_nodes=None, min_impurity_decrease=0.0,
min_impurity_split=None, min_samples_leaf=2,
min_samples_split=18, min_weight_fraction_leaf=0.0,
n_estimators=93, n_jobs=None, oob_score=True, random_state=25,
verbose=0, warm_start=False)
RandBestParams2= {'max_depth': 3, 'max_features': 3, 'min_samples_leaf': 2, 'min_samples_split
RandBestScore2= 0.24859766754045107

```

3.3 (d) Use sklearn to fit a support vector machines classifier model of cultivar with a gaussian radial basis function kernel on the four features used in parts b and c.

```
In [19]: svc= SVC(kernel="rbf", degree=2).fit(Xvars, yvals)

In [20]: param_dist3 = {'C':sp_uniform(loc=0.1,scale=10.0),
                        'gamma':["scale","auto"],
                        'shrinking':[True,False]}

random_search3 = RandomizedSearchCV(svc, param_distributions=param_dist3,
                                    n_iter=200, n_jobs=-1, cv=5, random_state=25,
                                    scoring='neg_mean_squared_error')

random_search3.fit(Xvars,yvals)
print('RandBestEstimator2=', random_search3.best_estimator_)
print('RandBestParams2=', random_search3.best_params_)
print('RandBestScore2=', -random_search3.best_score_)

RandBestEstimator2= SVC(C=3.3605112613782553, cache_size=200, class_weight=None, coef0=0.0,
                        decision_function_shape='ovr', degree=2, gamma='scale', kernel='rbf',
                        max_iter=-1, probability=False, random_state=None, shrinking=True,
                        tol=0.001, verbose=False)
RandBestParams2= {'C': 3.3605112613782553, 'gamma': 'scale', 'shrinking': True}
RandBestScore2= 0.14772727272727273
```

3.4 (e) Use neural net MLPClassifier to fit a multiple hidden layer neural network model of cultivar.

```
In [21]: mlp= MLPClassifier(activation='tanh', solver='lbfgs', alpha=0.1).fit(Xvars, yvals)

In [23]: param_dist4 = {'hidden_layer_sizes':sp_randint(1, 100),
                        'activation':["logistic", "relu"],
                        'alpha': sp_uniform(0.1, 10.0)}

random_search4 = RandomizedSearchCV(mlp, param_distributions=param_dist4,
                                    n_iter=200, n_jobs=-1, cv=5, random_state=25,
                                    scoring='neg_mean_squared_error')

random_search4.fit(Xvars,yvals)
print('RandBestEstimator2=', random_search4.best_estimator_)
print('RandBestParams2=', random_search4.best_params_)
print('RandBestScore2=', -random_search4.best_score_)

RandBestEstimator2= MLPClassifier(activation='relu', alpha=1.0882422675126213, batch_size='auto',
                                  beta_1=0.9, beta_2=0.999, early_stopping=False, epsilon=1e-08,
                                  hidden_layer_sizes=24, learning_rate='constant',
                                  learning_rate_init=0.001, max_iter=200, momentum=0.9,
                                  n_iter_no_change=10, nesterovs_momentum=True, power_t=0.5,
```

```
random_state=None, shuffle=True, solver='lbfgs', tol=0.0001,  
validation_fraction=0.1, verbose=False, warm_start=False)  
RandBestParams2= {'activation': 'relu', 'alpha': 1.0882422675126213, 'hidden_layer_sizes': 24}  
RandBestScore2= 0.07954545454545454
```

3.5 (f) Which of the above three models do you think is the best predictor of cultivar?

In terms of the MSE value, neural net MSLP is the best with the lowest MSE.

In []: